



## APPLICATION OF MACHINE LEARNING IN INDUSTRIES LAB

Name: Nidhi Mishra

SAP ID: 500097158

Roll no: R2142211276

Batch: B6

SUBMISSION TO:

**Dr. Pallabi Sharma**

Assistant Professor

School of Computer Science

UPES, Dehradun

<b><u>S.No.</u></b>	<b><u>Title</u></b>	<b><u>Page No.</u></b>
1	Introduction to Pandas and Numpy.	3-10
2	Wine quality prediction.	11-14
3	House Price prediction.	15-16
4	Air Quality Prediction.	17-20
5	Credit Card fault prediction.	21-23
6	Hypothesis Testing.	24-26
7	Neural Network (LSTM) model on sequential dataset.	27-28
8	Compare the LSTM (Neural network model) & CNN (Deep learning model).	29-31

## **EXPERIMENT-1**

### **Introduction to Pandas and Numpy**

#### **CODE-**

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

#### **# Task 1: Basic DataFrame Operations**

```
df = pd.read_csv('https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data',
                 header=None,
                 names=['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
                       'class'])
```

```
print(df.head())
print("\nMissing values:")
print(df.isnull().sum())
```

```
df.dropna(inplace=True)
```

```
print("\nSummary of the dataset:")
print(df.describe())
```

```
subset_label = df[['sepal_length', 'sepal_width', 'class']]
subset_position = df.iloc[:, [0, 1, 4]]
conditioned_df = df[df['sepal_length'] > 5.0]
```

### **# Task 2: Data Cleaning and Preprocessing**

```
missing_values = df.isnull().sum()  
df.fillna(df.mean(), inplace=True
```

```
df['class_encoded'] = df['class'].astype('category').cat.codes
```

```
grouped_data = df.groupby('class')
```

```
aggregated_results = grouped_data.agg({'sepal_length': ['sum',  
'ean', 'count']})
```

### **# Task 3: Merge two different datasets using different types of joins**

```
df1 = pd.read_csv('dataset1.csv')
df2 = pd.read_csv('dataset2.csv')

inner_join = pd.merge(df1, df2, how='inner',
on='common_column')
outer_join = pd.merge(df1, df2, how='outer',
on='common_column')
left_join = pd.merge(df1, df2, how='left', on='common_column')
right_join = pd.merge(df1, df2, how='right',
on='common_column')
```

### **# Task 4: Visualization**

```
df.plot.bar(x='class', y='sepal_length', title='Bar Plot')
df.plot.line(x='sepal_length', y='petal_length', title='Line Plot')
df.plot.scatter(x='sepal_length', y='petal_length', title='Scatter
Plot')
```

```
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
```

```
df.hist()
df.boxplot()
```

### **# Task 5: Basic NumPy Operations**

```
arr = np.arange(1, 11)
arr2 = np.arange(11, 21)
```

```
add_result = arr + arr2  
subtract_result = arr - arr2
```

```
multiply_result = arr * arr2  
divide_result = arr / arr2
```

### **# Task 6: Array Manipulation**

```
arr_reshaped = arr.reshape(2, 5)  
arr_transposed = arr_reshaped.T  
arr_flattened = arr_transposed.flatten()  
stacked_arrays = np.vstack((arr, arr2))
```

### **# Task 7: Statistical Operations**

```
mean_value = np.mean(arr)  
median_value = np.median(arr)  
std_deviation = np.std(arr)  
max_value = np.max(arr)  
min_value = np.min(arr)  
normalized_arr = (arr - mean_value) / std_deviation
```

### **# Task 8: Boolean Indexing**

```
bool_arr = arr > 5  
filtered_arr = arr[bool_arr]
```

### **# Task 9: Random Module**

```
random_matrix = np.random.rand(3, 3)  
random_integers = np.random.randint(1, 100, 10)  
np.random.shuffle(arr)
```

### **# Task 10: Universal Functions (ufunc)**

```
sqrt_arr = np.sqrt(arr)  
exp_arr = np.exp(arr)
```

## **# Task 11: Linear Algebra Operations**

```
mat_a = np.random.rand(3, 3)
vec_b = np.random.rand(3, 1)
result = np.dot(mat_a, vec_b)
```

## **# Task 12: Broadcasting**

```
matrix = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
mean_per_row = matrix.mean(axis=1, keepdims=True)
result_broadcasting = matrix - mean_per_row

plt.show()
```



## OUTPUT-

	sepal_length	sepal_width	petal_length	petal_width	class
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

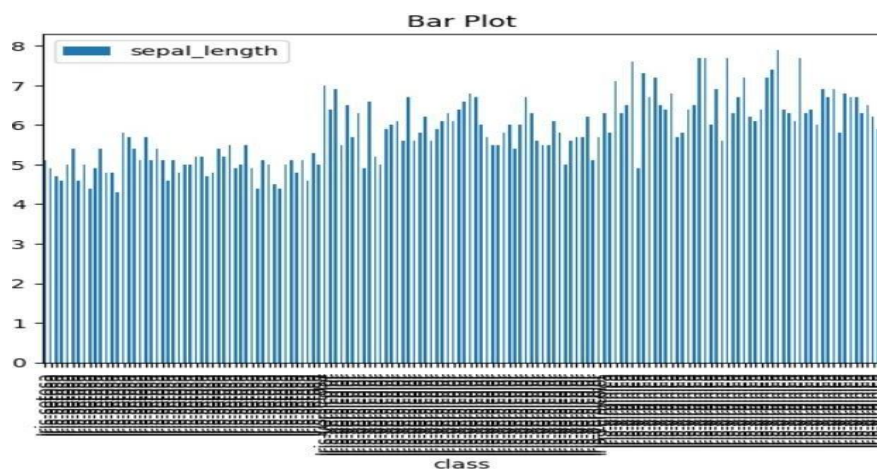
Missing values:

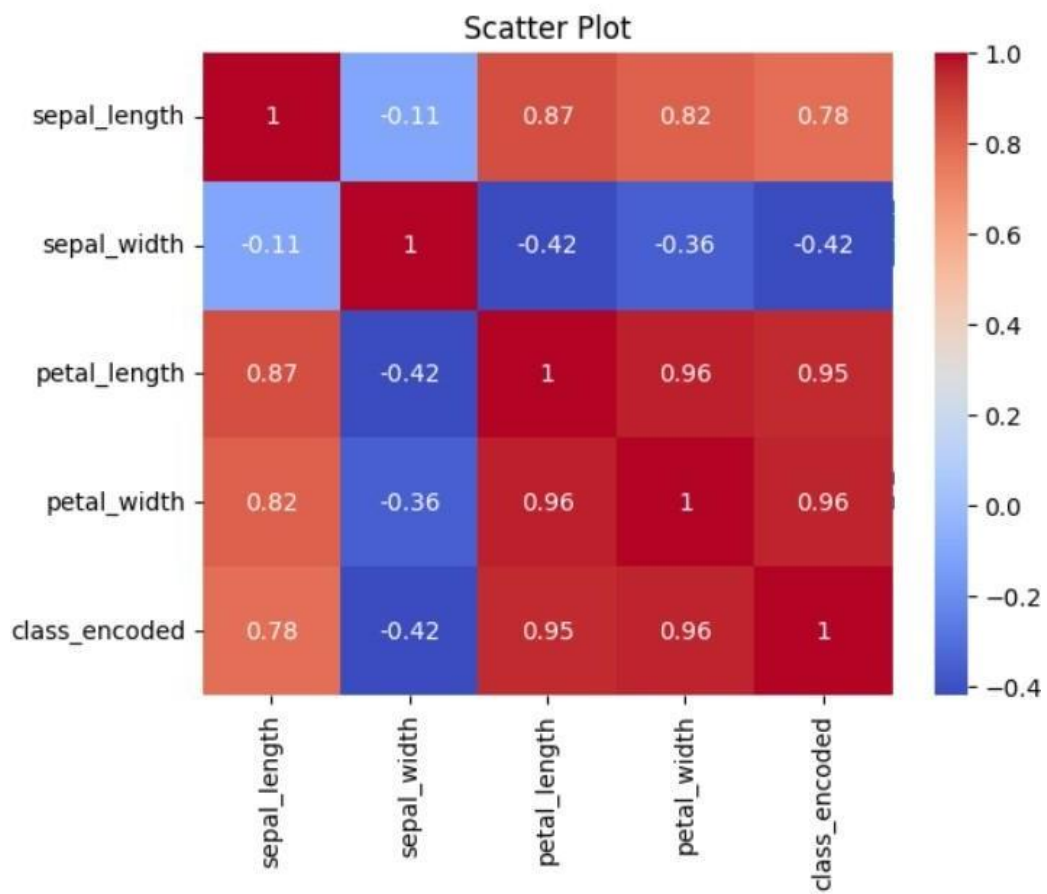
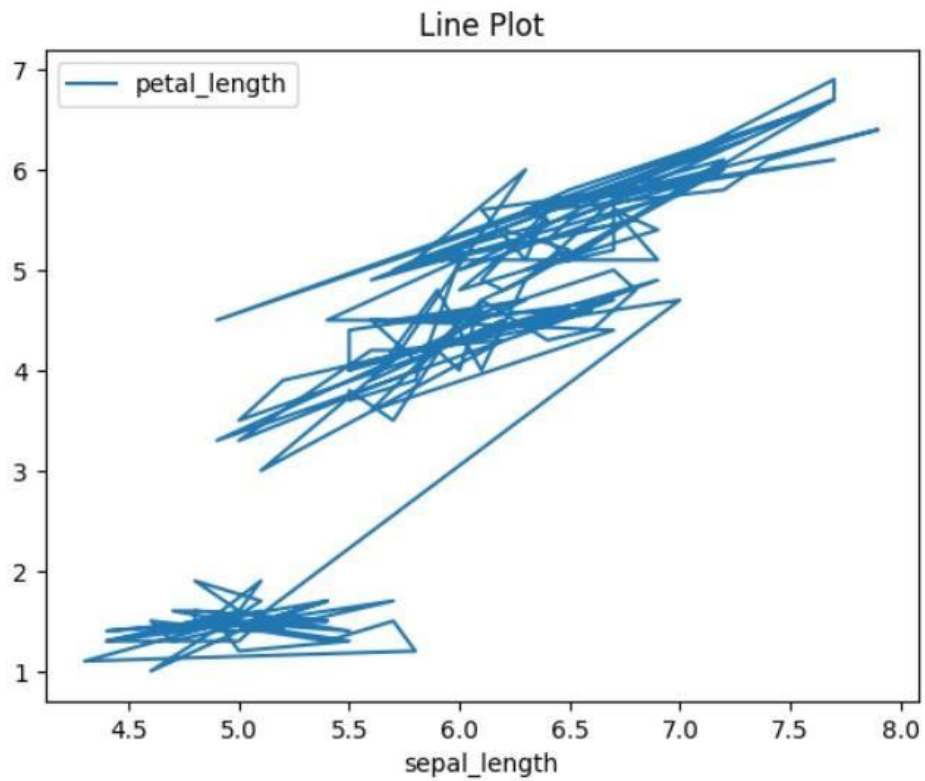
sepal_length	0
sepal_width	0
petal_length	0
petal_width	0
class	0

dtype: int64

Summary of the dataset:

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000





## **EXPERIMENT-2**

### **Wine quality prediction**

#### **CODE-**

```
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
import seaborn as sns
```

```
wine_data = pd.read_csv("C:/Users/pc/Downloads/archive
(6)/WineQT.csv")
```

```
# Split features and target variable
X = wine_data.drop('quality', axis=1)
y = wine_data['quality']
```

```
# Splitting the dataset into the Training set and Test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

```
# Feature scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
# Initialize the Random Forest classifier
rf_classifier = RandomForestClassifier(random_state=42)
```

```
# Define hyperparameters for tuning
param_grid = {
    'n_estimators': [50, 100, 150],
```

```

    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Hyperparameter tuning using GridSearchCV
grid_search = GridSearchCV(estimator=rf_classifier,
param_grid=param_grid, cv=5)
grid_search.fit(X_train_scaled, y_train)

# Best parameters found
best_params = grid_search.best_params_
print("Best Parameters:", best_params)

# Train the classifier with best parameters
best_rf_classifier = RandomForestClassifier(**best_params,
random_state=42)
best_rf_classifier.fit(X_train_scaled, y_train)

# Predict the test set results
y_pred = best_rf_classifier.predict(X_test_scaled)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# Plotting feature importances
feature_importances = best_rf_classifier.feature_importances_
feature_names = X.columns
sorted_idx = feature_importances.argsort()

plt.figure(figsize=(10, 6))

```

```

plt.barh(range(len(sorted_idx)), feature_importances[sorted_idx],
align='center')
plt.yticks(range(len(sorted_idx)), [feature_names[i] for i in
sorted_idx])
plt.xlabel('Feature Importance')
plt.title('Random Forest Feature Importance')
plt.show()

# Plotting confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.title('Confusion Matrix')
plt.show()

```



## OUTPUT-

	precision	recall	f1-score	support
4	0.00	0.00	0.00	6
5	0.73	0.75	0.74	96
6	0.64	0.71	0.67	99
7	0.76	0.62	0.68	26
8	0.00	0.00	0.00	2
accuracy			0.69	229
macro avg	0.43	0.41	0.42	229
weighted avg	0.67	0.69	0.68	229

```

Best Parameters: {'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
Accuracy: 0.6899563318777293

```



## **EXPERIMENT-3**

### **House Price prediction**

#### **CODE-**

```
import pandas as pd
from sklearn.preprocessing import OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_percentage_error

# Read the dataset
dataset = pd.read_excel("C:/Users/pc/Downloads/HousePricePrediction.xlsx")

# Drop 'Id' column and handle missing values
dataset.drop(['Id'], axis=1, inplace=True)
new_dataset = dataset.dropna()

# Identify categorical, integer, and float variables
object_cols = new_dataset.select_dtypes(include=['object']).columns.tolist()

# Perform one-hot encoding for categorical variables
OH_encoder = OneHotEncoder(sparse=False)
OH_cols = pd.DataFrame(OH_encoder.fit_transform(new_dataset[object_cols]))
OH_cols.index = new_dataset.index
OH_cols.columns = OH_encoder.get_feature_names_out(object_cols)
df_final = new_dataset.drop(object_cols, axis=1)
df_final = pd.concat([df_final, OH_cols], axis=1)

# Split dataset into features and target variable
X = df_final.drop(['SalePrice'], axis=1)
Y = df_final['SalePrice']
```

```
# Split the dataset into training and validation sets
X_train, X_valid, Y_train, Y_valid = train_test_split(X, Y,
train_size=0.8, test_size=0.2, random_state=0)
```

```
# Train Support Vector Regression model
model_SVR = SVR()
model_SVR.fit(X_train, Y_train)
Y_pred_SVR = model_SVR.predict(X_valid)
print("MAPE for SVR:", mean_absolute_percentage_error(Y_valid,
Y_pred_SVR))
```

```
# Train Random Forest Regression model
model_RFR = RandomForestRegressor(n_estimators=10)
model_RFR.fit(X_train, Y_train)
Y_pred_RFR = model_RFR.predict(X_valid)
print("MAPE for Random Forest:",
mean_absolute_percentage_error(Y_valid, Y_pred_RFR))
```

```
# Train Linear Regression model
model_LR = LinearRegression()
model_LR.fit(X_train, Y_train)
Y_pred_LR = model_LR.predict(X_valid)
print("MAPE for Linear Regression:",
mean_absolute_percentage_error(Y_valid, Y_pred_LR))
```

## OUTPUT-

```
PS C:\Users\pc> & C:/Users/pc/AppData/Local/Programs/Python/Python311/python.exe "c:/Users/pc/OneDrive/Desktop/ML in Industries/hpp.py"
C:\Users\pc\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\preprocessing\_encoders.py:972: FutureWarning: 'sparse' was renamed to 'sparse_output'
in version 1.2 and will be removed in 1.4. 'sparse_output' is ignored unless you leave 'sparse' to its default value.
  warnings.warn(
MAPE for SVR: 0.3009689871130911
MAPE for Random Forest: 0.14539550204661414
MAPE for Linear Regression: 0.2064918635154824
```



## **EXPERIMENT-4**

### **Air Quality Prediction**

#### **CODE-**

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

# Load the dataset
city_day = pd.read_csv("C:/Users/pc/Downloads/archive
(8)/city_day.csv")

# Drop rows with missing values
cleaned_data = city_day.dropna()

# Determine the number of unique cities in the dataset
no_of_cities = len(cleaned_data.City.unique())
print("Number of unique cities:", no_of_cities)

# Determine the unique AQI buckets in the dataset
aqi_buckets = cleaned_data.AQI_Bucket.unique()
print("Unique AQI Buckets:", aqi_buckets)

# Filter data for each AQI bucket
moderate = cleaned_data.AQI[cleaned_data.AQI_Bucket ==
'Moderate']
poor = cleaned_data.AQI[cleaned_data.AQI_Bucket == 'Poor']
very_poor = cleaned_data.AQI[cleaned_data.AQI_Bucket == 'Very
Poor']
satisfactory = cleaned_data.AQI[cleaned_data.AQI_Bucket ==
'Satisfactory']
good = cleaned_data.AQI[cleaned_data.AQI_Bucket == 'Good']
severe = cleaned_data.AQI[cleaned_data.AQI_Bucket == 'Severe']
```

```

# Selecting numerical columns for correlation analysis
gasses = cleaned_data.select_dtypes(include=np.float64)

# Calculating correlation between gases and AQI
corr = gasses.corr().AQI

# Identifying columns to drop based on correlation threshold
col_to_drop = corr[abs(corr) < 0.45].index

# Dropping columns with low correlation
gasses = gasses.drop(columns=col_to_drop)

# Defining target variable (AQI) and features (gases)
y = gasses.AQI
X = gasses.drop(columns='AQI')

# Normalize target variable (AQI)
y_max = y.max()
y = y / y_max

# Normalize features
X_max = X.max()
X = X / X_max

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Display model's performance

```

```
print("\nModel Evaluation:")
print("Mean Squared Error:", mse)
print("R-squared:", r2)



# Evaluate the model using cross-validation
cv_scores = cross_val_score(model, X_train, y_train, cv=5)
print("Cross-Validation Scores:", cv_scores)
print("Mean CV Score:", np.mean(cv_scores))

# Visualize model predictions vs actual air quality measurements
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred)
plt.xlabel("Actual AQI")
plt.ylabel("Predicted AQI")
plt.title("Model Predictions vs Actual AQI")
plt.show()

# Interpret trained model to understand feature importance
coefficients = pd.Series(model.coef_, index=X.columns)
print("Feature Importance:\n", coefficients)

# Visualize feature importances using bar plot
coefficients.plot(kind='barh')
plt.xlabel("Coefficient Value")
plt.ylabel("Feature")
plt.title("Feature Importance")
plt.show()
```

## OUTPUT-



## **EXPERIMENT-5**

### **Credit Card fault prediction**

#### **CODE-**

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import gridspec

data = pd.read_csv("C:/Users/pc/Downloads/archive
(12)/creditcard.csv")

data.head()

print(data.shape)
print(data.describe())

fraud = data[data['Class'] == 1]
valid = data[data['Class'] == 0]
outlierFraction = len(fraud)/float(len(valid))
print(outlierFraction)
print('Fraud Cases: {}'.format(len(data[data['Class'] == 1])))
print('Valid Transactions: {}'.format(len(data[data['Class'] == 0])))
print("Amount details of the fraudulent transaction")
fraud.Amount.describe()
print("details of valid transaction")
valid.Amount.describe()

corrmat = data.corr()
fig = plt.figure(figsize = (12, 9))
sns.heatmap(corrmat, vmax = .8, square = True)
plt.show()

X = data.drop(['Class'], axis = 1)
Y = data["Class"]
```

```
print(X.shape)
print(Y.shape)

xData = X.values
yData = Y.values

from sklearn.model_selection import train_test_split

xTrain, xTest, yTrain, yTest = train_test_split(
    xData, yData, test_size = 0.2, random_state = 42)

from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier()
rfc.fit(xTrain, yTrain)

yPred = rfc.predict(xTest)

from sklearn.metrics import classification_report, accuracy_score
from sklearn.metrics import precision_score, recall_score
from sklearn.metrics import f1_score, matthews_corrcoef
from sklearn.metrics import confusion_matrix

n_outliers = len(fraud)
n_errors = (yPred != yTest).sum()
print("The model used is Random Forest classifier")

acc = accuracy_score(yTest, yPred)
print("The accuracy is {}".format(acc))

prec = precision_score(yTest, yPred)
print("The precision is {}".format(prec))

rec = recall_score(yTest, yPred)
print("The recall is {}".format(rec))

f1 = f1_score(yTest, yPred)
```

```
print("The F1-Score is {}".format(f1))
```

```
MCC = matthews_corrcoef(yTest, yPred)
```

```
print("The Matthews correlation coefficient is {}".format(MCC))
```

```
LABELS = ['Normal', 'Fraud']
```

```
conf_matrix = confusion_matrix(yTest, yPred)
```

```
plt.figure(figsize=(12, 12))
```

```
sns.heatmap(conf_matrix, xticklabels=LABELS,  
            yticklabels=LABELS, annot=True, fmt="d");
```

```
plt.title("Confusion matrix")
```

```
plt.ylabel("True class")
```

```
plt.xlabel("Predicted class")
```


```
plt.show()
```

**Output-**

```
PS C:\Users\pc> & C:/Users/pc/AppData/Local/Programs/Python/Python311/python.exe "c:/Users/pc/OneDrive/Desktop/ML in Industries/credit_card_fraud/
detection.py"
(284807, 31)

   count  Time      V1      V2      V3      ...      V27      V28      Amount      Class
mean  284807.000000  2.848070e+05  2.848070e+05  2.848070e+05  ...  2.848070e+05  2.848070e+05  284807.000000  284807.000000
std    94813.859575  1.168375e-15  3.416908e-16  -1.379537e-15  ... -3.660091e-16  -1.227390e-16  88.349619  0.001727
std    47488.145955  1.958696e+00  1.651309e+00  1.516255e+00  ...  4.036325e-01  3.300833e-01  250.120109  0.041527
min     0.000000  -5.640751e+01  -7.271573e+01  -4.832559e+01  ... -2.256568e+01  -1.543008e+01  0.000000  0.000000
25%    54201.500000  -9.203734e-01  -5.985499e-01  -8.903648e-01  ... -7.083953e-02  -5.295979e-02  5.600000  0.000000
50%    84692.000000  1.810880e-02  6.548556e-02  1.798463e-01  ...  1.342146e-03  1.124383e-02  22.000000  0.000000
75%   139320.500000  1.315642e+00  8.037239e-01  1.027196e+00  ...  9.104512e-02  7.827995e-02  77.165000  0.000000
max   172792.000000  2.454930e+00  2.205773e+01  9.382558e+00  ...  3.161220e+01  3.384781e+01  25691.160000  1.000000

[8 rows x 31 columns]
0.0017304750013189597
Fraud Cases: 492
Valid Transactions: 284315
Amount details of the fraudulent transaction
details of valid transaction
```



## **EXPERIMENT-6**

### **HYPOTHESIS TESTING**

#### **CODE-**

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats

np.random.seed(0)
group1 = np.random.normal(loc=10, scale=2, size=100)
group2 = np.random.normal(loc=12, scale=2, size=100)
group3 = np.random.normal(loc=15, scale=2, size=100)

data = pd.DataFrame({
    'Group1': group1,
    'Group2': group2,
    'Group3': group3
})

sns.boxplot(data=data)
plt.title('Boxplot of the Random Dataset')
plt.xlabel('Groups')
```



```
plt.ylabel('Values')  
plt.show()
```

```
t_statistic, p_value = stats.ttest_ind(group1, group2)  
print("T-Test between Group1 and Group2:")  
print("T-statistic:", t_statistic)  
print("P-value:", p_value)
```

```
z_statistic, p_value = stats.zscore([group1, group2])  
print("\nZ-Test between Group1 and Group2:")  
print("Z-statistic:", z_statistic[0] - z_statistic[1])  
print("P-value:", p_value)
```

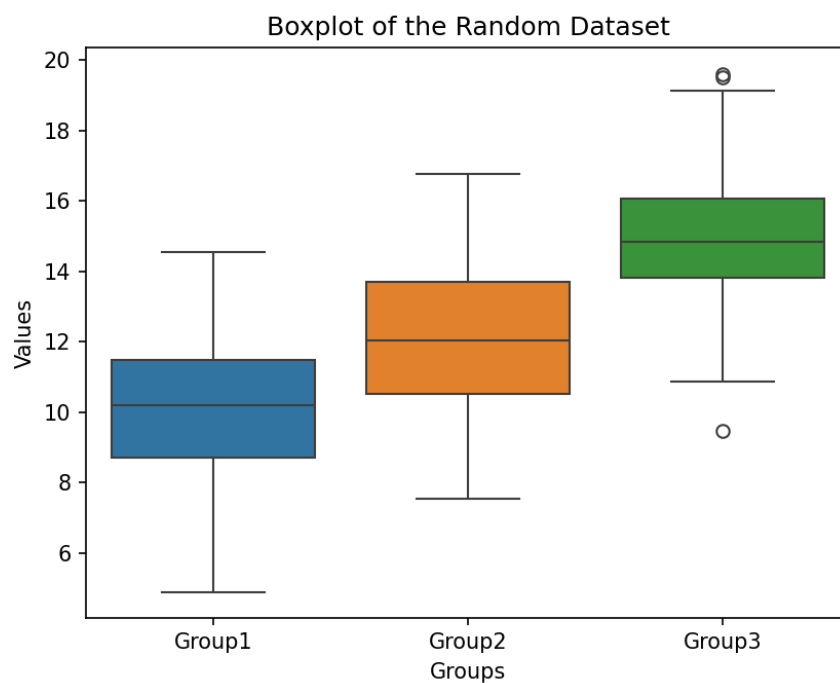
```
f_statistic, p_value = stats.f_oneway(group1, group2, group3)  
print("\nANOVA:")  
print("F-statistic:", f_statistic)  
print("P-value:", p_value)
```

## OUTPUT-

```
PS C:\Users\pc> & C:/Users/pc/AppData/Local/Programs/Python/Python311/python.exe "c:/Users/pc/OneDrive/nova.py"
T-Test between Group1 and Group2:
T-statistic: -7.0414273690132765
P-value: 3.059820094513985e-11

Z-Test between Group1 and Group2:
Z-statistic: -2.0
P-value: [ 1. -1. -1. -1. -1.  1. -1.  1.  1.  1.  1.  1. -1.  1.  1.  1.  1.
  1.  1.  1. -1.  1.  1. -1.  1.  1.  1.  1. -1.  1.  1.  1.  1.  1.
  1. -1.  1.  1.  1.  1.  1. -1.  1.  1.  1.  1.  1.  1.  1.  1.  1.
  1. -1.  1. -1.  1.  1.  1.  1.  1.  1. -1.  1.  1. -1.  1.  1.  1.
  1.  1. -1.  1.  1.  1.  1.  1.  1. -1. -1.  1.  1. -1.  1.  1.
  1.  1.  1. -1.  1.  1.  1.  1.  1.  1.]

ANOVA:
F-statistic: 141.5873093832312
P-value: 6.544281608055403e-44
```



## **EXPERIMENT-7**

### **Neural network(LSTM) model on sequential dataset**

#### **Code-**

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import imdb
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding,
SpatialDropout1D
from tensorflow.keras.preprocessing.sequence import pad_sequences

max_features = 20000
maxlen = 100
(x_train, y_train), (x_test, y_test) =
imdb.load_data(num_words=max_features)

x_train = pad_sequences(x_train, maxlen=maxlen)
x_test = pad_sequences(x_test, maxlen=maxlen)

embedding_dim = 128
lstm_units = 64

model = Sequential([
    Embedding(max_features, embedding_dim, input_length=maxlen),
    SpatialDropout1D(0.2),
    LSTM(lstm_units, dropout=0.2, recurrent_dropout=0.2),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

batch_size = 128
epochs = 5
model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs,
validation_data=(x_test, y_test))
```

```
loss, accuracy = model.evaluate(x_test, y_test)
print("Test Loss:", loss)
print("Test Accuracy:", accuracy)
```

## Output-

---

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 [=====] - 0s 0us/step
Epoch 1/5
196/196 [=====] - 97s 478ms/step - loss: 0.4498 - accuracy: 0.7788 - val_loss: 0.3444 - val_accuracy: 0.8502
Epoch 2/5
196/196 [=====] - 103s 524ms/step - loss: 0.2531 - accuracy: 0.9016 - val_loss: 0.3918 - val_accuracy: 0.8467
Epoch 3/5
196/196 [=====] - 103s 528ms/step - loss: 0.1755 - accuracy: 0.9350 - val_loss: 0.3758 - val_accuracy: 0.8427
Epoch 4/5
196/196 [=====] - 102s 519ms/step - loss: 0.1280 - accuracy: 0.9546 - val_loss: 0.4214 - val_accuracy: 0.8352
Epoch 5/5
196/196 [=====] - 92s 468ms/step - loss: 0.0942 - accuracy: 0.9672 - val_loss: 0.5009 - val_accuracy: 0.8285
782/782 [=====] - 24s 30ms/step - loss: 0.5009 - accuracy: 0.8285
Test Loss: 0.5008736848831177
Test Accuracy: 0.828519999809265
```

## **EXPERIMENT-8**

### **Compare the LSTM(Neural network model)&CNN(Deep learning model)**

#### **CODE-**

```
import numpy as np
import matplotlib.pyplot as plt
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D
from keras.utils import to_categorical

(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train / 255.0
X_test = X_test / 255.0

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)

def create_sequential_model():
    model = Sequential([
        Flatten(input_shape=(28, 28)),
        Dense(128, activation='relu'),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
    return model
```

```

sequential_model = create_sequential_model()
sequential_history = sequential_model.fit(X_train, y_train,
epochs=10, batch_size=128, validation_data=(X_test, y_test))

def create_cnn_model():
    model = Sequential([
        Conv2D(32, kernel_size=(3, 3), activation='relu',
input_shape=(28, 28, 1)),
        MaxPooling2D(pool_size=(2, 2)),
        Flatten(),
        Dense(128, activation='relu'),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
    return model

X_train_cnn = X_train.reshape(-1, 28, 28, 1)
X_test_cnn = X_test.reshape(-1, 28, 28, 1)

cnn_model = create_cnn_model()
cnn_history = cnn_model.fit(X_train_cnn, y_train, epochs=10,
batch_size=128, validation_data=(X_test_cnn, y_test))

def plot_history(history, title):
    plt.plot(history.history['accuracy'], label='accuracy')
    plt.plot(history.history['val_accuracy'], label='val_accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.title(title)
    plt.legend()

```

```
plt.show()
```

```
plot_history(sequential_history, 'Sequential Model')
```

```
plot_history(cnn_history, 'CNN Model')
```

## OUTPUT-

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 4s 0us/step
C:\Users\pc\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\layers\reshaping\flatten.py:37: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model instead.
  super().__init__(**kwargs)
2024-04-15 10:31:18.078682: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
Epoch 1/10
469/469 ————— 2s 2ms/step - accuracy: 0.8295 - loss: 0.6176 - val_accuracy: 0.9454 - val_loss: 0.1879
Epoch 2/10
469/469 ————— 1s 2ms/step - accuracy: 0.9506 - loss: 0.1744 - val_accuracy: 0.9605 - val_loss: 0.1311
Epoch 3/10
469/469 ————— 1s 2ms/step - accuracy: 0.9661 - loss: 0.1156 - val_accuracy: 0.9691 - val_loss: 0.1045
Epoch 4/10
469/469 ————— 1s 2ms/step - accuracy: 0.9749 - loss: 0.0879 - val_accuracy: 0.9705 - val_loss: 0.0958
Epoch 5/10
469/469 ————— 1s 2ms/step - accuracy: 0.9795 - loss: 0.0692 - val_accuracy: 0.9762 - val_loss: 0.0796
Epoch 6/10
469/469 ————— 1s 2ms/step - accuracy: 0.9835 - loss: 0.0559 - val_accuracy: 0.9742 - val_loss: 0.0797
Epoch 7/10
469/469 ————— 1s 2ms/step - accuracy: 0.9876 - loss: 0.0457 - val_accuracy: 0.9791 - val_loss: 0.0709
Epoch 8/10
469/469 ————— 1s 2ms/step - accuracy: 0.9890 - loss: 0.0396 - val_accuracy: 0.9788 - val_loss: 0.0718
Epoch 9/10
469/469 ————— 1s 2ms/step - accuracy: 0.9906 - loss: 0.0329 - val_accuracy: 0.9781 - val_loss: 0.0716
Epoch 10/10
469/469 ————— 1s 2ms/step - accuracy: 0.9930 - loss: 0.0277 - val_accuracy: 0.9765 - val_loss: 0.0703
```

```
C:\Users\pc\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:99: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in the model instead.
  super().__init__(
Epoch 1/10
469/469 ————— 7s 13ms/step - accuracy: 0.8736 - loss: 0.4448 - val_accuracy: 0.9763 - val_loss: 0.0766
Epoch 2/10
469/469 ————— 6s 13ms/step - accuracy: 0.9789 - loss: 0.0747 - val_accuracy: 0.9823 - val_loss: 0.0540
Epoch 3/10
469/469 ————— 6s 13ms/step - accuracy: 0.9862 - loss: 0.0454 - val_accuracy: 0.9820 - val_loss: 0.0491
Epoch 4/10
469/469 ————— 6s 13ms/step - accuracy: 0.9901 - loss: 0.0328 - val_accuracy: 0.9850 - val_loss: 0.0431
Epoch 5/10
469/469 ————— 6s 13ms/step - accuracy: 0.9929 - loss: 0.0238 - val_accuracy: 0.9861 - val_loss: 0.0401
Epoch 6/10
469/469 ————— 6s 13ms/step - accuracy: 0.9947 - loss: 0.0170 - val_accuracy: 0.9849 - val_loss: 0.0454
Epoch 7/10
469/469 ————— 6s 13ms/step - accuracy: 0.9963 - loss: 0.0147 - val_accuracy: 0.9845 - val_loss: 0.0438
Epoch 8/10
469/469 ————— 6s 13ms/step - accuracy: 0.9971 - loss: 0.0100 - val_accuracy: 0.9858 - val_loss: 0.0464
Epoch 9/10
469/469 ————— 6s 13ms/step - accuracy: 0.9978 - loss: 0.0081 - val_accuracy: 0.9846 - val_loss: 0.0479
Epoch 10/10
469/469 ————— 6s 13ms/step - accuracy: 0.9983 - loss: 0.0066 - val_accuracy: 0.9851 - val_loss: 0.0473
```